

```
In [82]: import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import confusion_matrix, accuracy_score, precision_score, recall_score
```

```
In [83]: df = pd.read_csv("Social_Network_Ads.csv")
df
```

Out[83]:

	User ID	Gender	Age	EstimatedSalary	Purchased
0	15624510	Male	19	19000	0
1	15810944	Male	35	20000	0
2	15668575	Female	26	43000	0
3	15603246	Female	27	57000	0
4	15804002	Male	19	76000	0
...	...	...	...	...	...
395	15691863	Female	46	41000	1
396	15706071	Male	51	23000	1
397	15654296	Female	50	20000	1
398	15755018	Male	36	33000	0
399	15594041	Female	49	36000	1

400 rows × 5 columns

```
In [84]: df["Gender"] = df["Gender"].replace({"Male": 0, "Female": 1})
df
C:\Users\Ritesh Kolt\AppData\Local\Temp\ipykernel_19336\3117893451.py:1: FutureWarning: Downcasting behavior in `replace` is deprecated and will be removed in a future version. To retain the old behavior, explicitly call `result.infer_objects(copy=False)`. To opt
-in to the future behavior, set `pd.set_option('future.no_silent_downcasting', True)`
df["Gender"] = df["Gender"].replace({"Male": 0, "Female": 1})
```

Out[84]:

	User ID	Gender	Age	EstimatedSalary	Purchased
0	15624510	0	19	19000	0
1	15810944	0	35	20000	0
2	15668575	1	26	43000	0
3	15603246	1	27	57000	0
4	15804002	0	19	76000	0
...	...	...	...	...	...
395	15691863	1	46	41000	1
396	15706071	0	51	23000	1
397	15654296	1	50	20000	1
398	15755018	0	36	33000	0
399	15594041	1	49	36000	1

400 rows × 5 columns

```
In [85]: df.columns
```

Out[85]: Index(['User ID', 'Gender', 'Age', 'EstimatedSalary', 'Purchased'], dtype='object')

```
In [86]: x = df[['User ID', 'Gender', 'Age', 'EstimatedSalary']]
y = df['Purchased']
```

```
In [87]: x
```

Out[87]:

	User ID	Gender	Age	EstimatedSalary
0	15624510	0	19	19000
1	15810944	0	35	20000
2	15668575	1	26	43000
3	15603246	1	27	57000
4	15804002	0	19	76000
...	...	...	...	...
395	15691863	1	46	41000
396	15706071	0	51	23000
397	15654296	1	50	20000
398	15755018	0	36	33000
399	15594041	1	49	36000

400 rows × 4 columns

```
In [88]: y
```

Out[88]:

0	0
1	0
2	0
3	0
4	0
...	..
395	1
396	1
397	1
398	0
399	1

Name: Purchased, Length: 400, dtype: int64

```
In [89]: x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.25,random_state=29)
```

```
In [90]: model = LogisticRegression()
model.fit(x_train,y_train)
```

Out[90]:

▼ LogisticRegression

LogisticRegression()

```
In [91]: y_pred = model.predict(x_test)
```

```
In [92]: y_pred
```

Out[92]:

```
array([[0, 0, 0, 0, 0, 0, 1, 1, 1, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
        0, 0, 1, 1, 0, 1, 0, 0, 1, 0, 1, 0, 0, 0, 0, 1, 0, 1, 0, 1, 0, 0,
        0, 0, 0, 0, 1, 1, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0,
        0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0,
        1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0,
        1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 1])
```

```
In [93]: model.score(x_train,y_train)
```

Out[93]: 0.8633333333333333

```
In [94]: model.score(x,y)
```

Out[94]: 0.85

```
In [96]: cm = confusion_matrix(y_test, y_pred)
cm
```

Out[96]:

```
array([[63,  6],
       [13, 18]])
```

```
In [99]: tn, fp, fn, tp = confusion_matrix(y_test, y_pred).ravel()
```

```
In [100]: print(tn,fp,fn,tp)

63 6 13 18
```

```
In [101]: a = accuracy_score(y_test,y_pred)
a
```

Out[101]: 0.81

```
In [103]: e = 1 - a      #error value
e
```

Out[103]: 0.18999999999999995

```
In [104]: precision_score(y_test,y_pred)
```

Out[104]: 0.75

```
In [105]: recall_score(y_test,y_pred)
```

